

# Jinning's Tweet Search System

The search system is done in Java, as I'm more familiar with that. (Note that I have some experience with Python.)

The whole search problem is very similar to "searching words (tokens) in tweets". I used six different algorithms with different pros and cons: *Naive*, *TokenizedSearch*, *SimpleSuffixTree*, *MultiPartitionSuffixTree*, *SimpleTokenizedSuffixTree*, and *MultiPartitionTokenizedSuffixTree*.

## Note Before

- A. I did not only store the latest tweet as using that means we can not tune the ranking algorithm.
- B. To run the tests, use "run\_tests.sh".

## The Highest Performance Solution

The highest performance solution I can think of is to use a HashMap and put all tokens in all tweets into the map along with the Tweet indices. -- The given problem size is not very large, so I didn't partition the HashMaps. However, if we have too many tweets and tokens, we can partition the tokens and put them on different machines, and then use some well-studied reduce algorithms to get the result.

## Trading Some Performance for Better Search Result

Hashing everything sounds good, but what if a user:

- a) Searches "*BestTaco*" and lots of tweets contain "*TheBestTaco*"?
- b) Searches "*apple laptop*" and we return something with "*Have an apple in the laptop bag*" rather than "*The latest Apple laptop is so cool*"?

In these cases, we might want to change our search algorithm a bit and enable it to match substrings of the tokens or even substrings of the tweets. One simple (i.e. implementable in a few hundred lines without using third-party libs) way is to use **suffix tree**, and that's what I choose.

[Suffix tree](#) basically means that we put every suffix of a string into a trie tree, and do matching on the tree. In my implementation, I stored all twitter in a giant suffix tree, and mark the suffix with the tweet indices stored in a set in the tree node.

## Performance Concerns

A suffix tree is based on a trie tree, which is known to take a huge amount of memory if the string is not short. Luckily, the tweets are usually short. So the trie tree can be smaller. But still, more than 32GB of memory will be used when we're dealing with a large ( $> 100,000$ ) number of tweets.

For my suffix tree implementation, a big issue is that my code will use up a huge amount of memory when building the tree, as we will create and link the nodes per character. Also, Java will use tons of memory "by design". I used these tricks to help reduce memory consumption:

### Multi Partitioning

As a suffix tree's memory footprint is roughly ***quadratic in the length of the string*** and ***linear in the number of the string***. But in practice, this is only a rough upper bound. If we log the memory consumption of our suffix tree, we can observe that once we've built the tree, the memory footprint will drastically decrease as we will trim the tree down to its most optional form.

If we have more time, we can use some fancy algorithm to build the tree. Like recently [a parallel tree-building algorithm with  \$O\(N\)\$  time and  \$O\(\log^2 N\)\$  span](#) has been proposed.

### Optimization 1

As said before, I don't have time to thoroughly read and implement those algorithms. So I just used a simpler approach: Partition the tweets and put them in different suffix trees.

After some experiments, I used 20000 tweets per suffix tree if we have less than 60000 tweets in total; 10000 tweets per tree if we have less than 120000 tweets; and 5000 tweets if we have more than that.

With these settings, I compared the partitioned suffix tree's performance with naive implementation:

Problem size	Naive Time (ms)	Suffix Tree time (ms)	Speedup
Small	20706	70	296x
Medium	95606	658	145x
Large	184869	2736	68x

Algorithm "*MultiPartitionSuffixTree*", "*MultiPartitionTokenizedSuffixTree*" used this optimization.

## Optimization 2

Also, there's another way to do it if we don't mind scenario b mentioned above, which is to break tweets into tokens, and put the tokens into the suffix tree.

Algorithm "*SimpleTokenizedSuffixTree*", "*MultiPartitionTokenizedSuffixTree*" uses this optimization.

## The Sorting algorithm

I've implemented a "Recommender" class and an "AlgorithmHandler" interface for sorting the results. Here are the implemented algorithms:

- *Latest*
- *HighestViews*
- *HighestLikes*
- *HighestRetweets*
- *MyAwesomeModel*


The last algorithm is just a model considering the tweet's all stats, with a strong preference for the latest tweets, calculated as:

$$\begin{aligned} \text{match\_score} &= \text{views} + \text{likes} * 50 + \text{retweets} * 100 * \text{deteration\_factor}, \text{ where} \\ \text{deteration\_factor} &= (\text{LatestTime} + 1) / (\text{LatestTime} + 1 - \text{TweetTime}) \end{aligned}$$

## The Overall Performance

See charts for the overall performance:

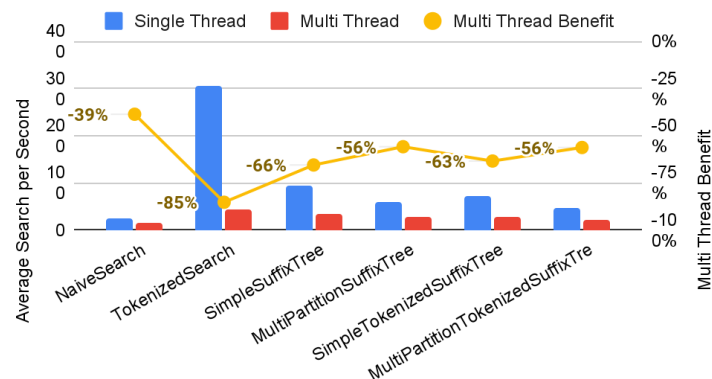
## 90 Percentile Performance of Different Algorithms



There's a weird drop for suffix trees when the number of tokens go from 2 to 3. I don't have time to look at this in depth. **If I have more time I plan on using something like a IntelliJ's Java profiler to narrow down why this happens.**

I also tried multi-threading the queries, but it seems that it can't bring any performance benefit on my machine. Probably due to the query is memory limited and multi-threading will just saturate the CPU -> DRAM path. If I have more time, I'll do some in-depth profiling using VTune to verify my theory:

## 90 Percentile Performance of Different Algorithms



# Extensions I Would Like To Have

1. Implement the spell corrector.
2. Cache the search results.
3. Separate the control and data in unit tests.
4. Currently, we have bubbles in `MultiThreadSearch` and `MultiThreadReduce`.

*MultiThreadSearch -> Join -> MultiThreadReduce*

We can make it to:

*MultiThreadSearch -> MultiThreadReduce*

5. Create a `SearchAbleFactory` for cleaner code.  
But it's painful to write Factory with generics in Java, and the code is not likely to live long, so... :)
6. Do in-depth studies for different search algorithms' performance, find optimization opportunities, and use the most optimal performance at different problem sizes.

# Appendix

## UML Diagram

